

DocuChat: Train Your AI on PDFs

Abstract

DocuChat is a Langchain-based Multi-PDF Chatbot that uses the RAG (Retrieval-Augmented Generation) framework to enable natural language interaction with multiple PDF documents. By integrating Ollama and Huggingface language models, it allows users to upload PDFs, ask questions, and receive contextually accurate responses. The system extracts text, generates embeddings, and stores them in a vector database for efficient retrieval, which are then used by language models to generate coherent answers. Designed for researchers, students, and professionals, DocuChat combines retrieval and generation to provide a scalable, user-friendly solution for analyzing and extracting insights from large documents.

Introduction

In the modern era of information overload, PDF documents have become one of the most widely used formats for storing and sharing critical information. From research papers and legal contracts to financial reports and educational materials, PDFs are ubiquitous across industries. However, extracting meaningful insights from these documents can be a daunting task. Traditional methods of manually searching through PDFs are not only time-consuming but also inefficient, especially when dealing with large volumes of text. This challenge is further compounded when users need to interact with multiple PDFs simultaneously, as maintaining context across documents becomes increasingly difficult.

To address these challenges, we present DocuChat, a Langchain-based application designed to revolutionize the way users interact with PDF documents. DocuChat leverages the RAG (Retrieval-Augmented Generation) framework, along with Ollama and Huggingface language models, to provide a natural language interface for querying, summarizing, and analyzing multiple PDF files. By combining advanced natural language processing (NLP) techniques with a user-friendly interface, DocuChat empowers users to extract valuable insights from their documents with ease. The retrieval phase involves extracting text, generating embeddings, and storing them in a vector database, while the generation phase uses language models to produce accurate, context-aware responses.

The core innovation of DocuChat lies in its ability to process and understand the content of multiple PDFs simultaneously using the RAG framework. Users can upload several PDF files, ask questions in natural language, and receive precise, context-aware responses. For instance, a researcher can upload multiple research papers and ask DocuChat to compare findings or summarize key points. Similarly, a student can upload textbooks and request explanations of complex topics. Professionals in legal, financial, or technical fields can use DocuChat to analyze contracts, reports, or manuals efficiently. By integrating retrieval and generation, DocuChat provides a scalable, efficient, and intuitive solution for interacting with PDF documents.

Objectives

The primary objective of this project is to develop a Langchain-based application that enables users to interact with multiple PDF files using natural language. The app will leverage Ollama and Huggingface language models to provide accurate, context-aware responses to user queries. Specific objectives include:

1. Building a system that can process and extract text from multiple PDF files.
2. Implementing a vector database for efficient storage and retrieval of document embeddings.
3. Integrating advanced language models for natural language understanding and response generation.
4. Creating a user-friendly interface for seamless interaction with the application.

Existing Systems

1. Traditional PDF Readers (e.g., Adobe Acrobat):
 - Basic keyword search; no natural language interaction or multi-PDF support.
2. PDF Text Extraction Tools (e.g., PyPDF2):
 - Extract raw text but lack querying or analysis capabilities.
3. Document Search Engines (e.g., Elasticsearch):
 - Keyword-based searches; no natural language understanding or summarization.
4. AI-Powered Tools (e.g., ChatPDF):
 - Limited to single PDFs; no context awareness across multiple documents.

DocuChat addresses these limitations by enabling natural language interaction with multiple PDFs, leveraging advanced NLP models for context-aware responses.

Problem Statement


The increasing reliance on PDF documents for storing critical information has created a need for efficient tools to interact with these files. Traditional methods of searching and extracting information from PDFs are manual, time-consuming, and often ineffective. Users require a solution that allows them to query multiple PDFs simultaneously and receive accurate, context-aware responses. This project addresses this problem by developing a Langchain-based application that leverages Ollama and Huggingface language models to provide a natural language interface for interacting with PDF documents.

Proposed System

DocuChat is a Langchain-based application that enables natural language interaction with multiple PDF files. It leverages Ollama and Huggingface language models to provide context-aware responses, summaries, and insights. Key features include:

1. Multi-PDF Support: Upload and query multiple PDFs simultaneously.
2. Advanced NLP: Uses Huggingface models for accurate, context-aware responses.
3. Efficient Retrieval: Embeds text chunks in a vector database for fast retrieval.

Architecture of the Model



System Requirements

Hardware Requirements:

1. Processor: Minimum Intel i5 12th gen or equivalent.
2. RAM: Minimum 8 GB (recommended: 16 GB or higher for handling large documents).
3. Storage: Minimum 5 GB free space (for storing embeddings and temporary files).

Software Requirements:

1. Operating System: Windows 10/11.
2. Libraries and Frameworks:
 - Langchain
 - Huggingface Transformers

- Ollama
 - PyPDF2 or PDFPlumber
 - Streamlit
 - Sentence Transformers (Huggingface)
3. Development Tools:
- IDE: VS Code, PyCharm, or Jupyter Notebook.
 - Version Control: Git.

Modules

1. PDF Upload and Text Extraction Module:
 - Handles PDF uploads and extracts text using PyPDF2 or PDFPlumber.
2. Text Chunking and Embedding Module:
 - Splits text into manageable chunks and generates embeddings using Huggingface models.
3. User Interface Module:
 - Provides a web-based interface for uploading PDFs and interacting with the chatbot.

Conclusion

The Langchain-based Multi-PDF Chatbot is a powerful tool that simplifies the process of interacting with PDF documents. By leveraging Ollama and Huggingface language models, the app provides accurate, context-aware responses to user queries. The system's modular architecture ensures scalability and efficiency, making it suitable for a wide range of applications, from research and education to business intelligence. This project demonstrates the potential of integrating advanced language models with user-friendly interfaces to solve real-world problems in document management and information retrieval.

References

1. Langchain Documentation: <https://www.langchain.com/>
2. Huggingface Transformers: <https://huggingface.co/transformers/>
3. Ollama GitHub Repository: <https://github.com/ollama/ollama>
4. PyPDF2 Documentation: <https://pypi.org/project/PyPDF2/>
5. FAISS Vector Database: <https://github.com/facebookresearch/faiss>
6. Streamlit Documentation: <https://docs.streamlit.io/>

DocuChat:

Train Your AI on PDFs

Submitted By,

H N Adith

U05AD22S0044

Dilan S

U05AD22S0050